

# AI Masterclass

Technical Generative AI Concepts Explained Simply



# Learning Journey Roadmap



## Goals

- ✓ Understand how GenAI technology works
- ✓ Feel comfortable exploring with GenAI tools
- ✓ Start applying GenAI technology safely and responsibly

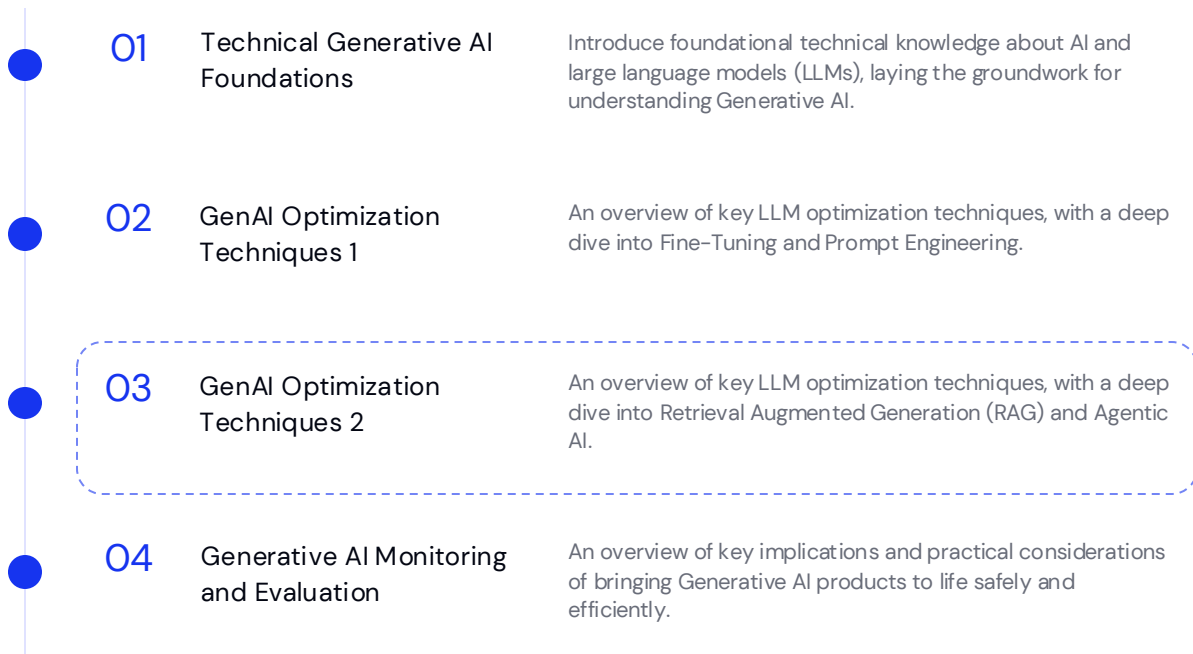
# Presented by



## Esteban D. Lopez

 Location	New York, NY
 Hometown	Quito, Ecuador
 Education	University of New Orleans – BS. Accounting Louisiana State University – MS. Accounting Columbia University – FinTech Bootcamp New York University – MS. AI and Machine Learning (Current)
 Professional Career	Hedge Funds Assurance Transformation & Economic Consulting AI & Data Technology Consulting AI Product Management

# Learning Journey Roadmap



## Goals


- ✓ Understand how GenAI technology works
- ✓ Feel comfortable exploring with GenAI tools
- ✓ Start applying GenAI technology safely and responsibly

# GenAI Optimization Techniques

## Part II



# Overview of Generative AI Principles



## Generative AI

AI that creates new content—such as text, code, images, or music—rather than just analyzing data.



## Neural Networks

Layers of connected artificial neurons that process data and learn complex patterns in large sets of data through training.



## Large Language Models

Very large neural networks trained on massive text datasets to generate human-like language.



## Transformer Architecture

The neural network design that powers modern LLMs using self-attention to understand word relationships.



## Natural Language Processing

The field of enabling computers to understand, interpret, and generate human language..



## Tokenization

The process of breaking text into small units (tokens) that an LLM can understand and process.,

# Optimization Techniques

01

## Fine-tuning

Training an LLM on custom data to specialize its behavior.

02

## Prompt Engineering

Designing effective prompts to guide model outputs.

03

## Retrieval-Augmented Generation (RAG)

Grounds LLMs with external knowledge sources.

04

## Agentic AI

Orchestrates LLMs as multi-step, tool-using agents with memory and reasoning.



## Cost

The amount of resource (data, compute, and engineering effort) needed to implement and maintain each technique.



## Implementation Efficiency

























How quickly and easily the technique can be deployed or iterated on in real-world workflows.



## Performance

The degree of improvement the technique delivers in output quality, accuracy, and reliability.

# Optimization Techniques

Technique	Description	Cost	Implementation
 <b>Fine-tuning</b>	Further training enhances model performance.	   	   
 <b>Retrieval Augmented Generation (RAG)</b>	Connects model to external data sources.	 	 
 <b>Prompt Engineering</b>	Refines questions to maximize response quality.		
 <b>Agentic AI</b>	Intelligent agents automate decisions and tasks.	  	  



## Part 2

### Retrieval-Augmented Generation:

Demonstrates how to inject recent or domain-specific knowledge into models without altering their weights, improving accuracy.

### Agentic AI Orchestration

Explains how to extend LLMs with reasoning, memory, and tools, enabling multi-step workflows and complex problem-solving.



# RAG (Retrieval Augmented Generation)

High-level overview of RAG process for LLMs.

01

## Retrieve Relevant Knowledge

- The system searches external sources (databases, documents, APIs) for context based on the user's query
- Uses a vector database and embeddings to find semantically similar content

02

## Augment the Prompt with Retrieved Context

- Injects the retrieved passages into the model's prompt
- Gives the model up-to-date and domain-specific information before generating

03

## Generate Answer Using the LLM

- The model produces an output that blends its own language skills with the retrieved facts
- This improves factual accuracy and reduces hallucinations

### RAG Pipeline Flow

#### User Asks Question

The user initiates the process by asking a question.

#### Retriever Scans Data

The retriever scans connected data sources for relevant information.

#### Relevant Documents Pulled

The retriever pulls in the relevant documents.

#### Generator Reads Documents

The generator reads the documents to understand the context.

#### Generator Creates Response

The generator creates a grounded response based on the documents.

# RAG – Analysis

## Benefits

RAG enhances accuracy and freshness by giving models real-time access to external knowledge without retraining.

## Drawbacks

It adds infrastructure complexity and depends on the quality of the retrieval data and indexing, which can affect speed and accuracy.

## Sample Use Cases

### Customer Support Bots

Providing real-time answers from product manuals and FAQs to reduce hallucinations.


### Enterprise Knowledge Search

Answering employee questions from internal documents without fine-tuning the model.

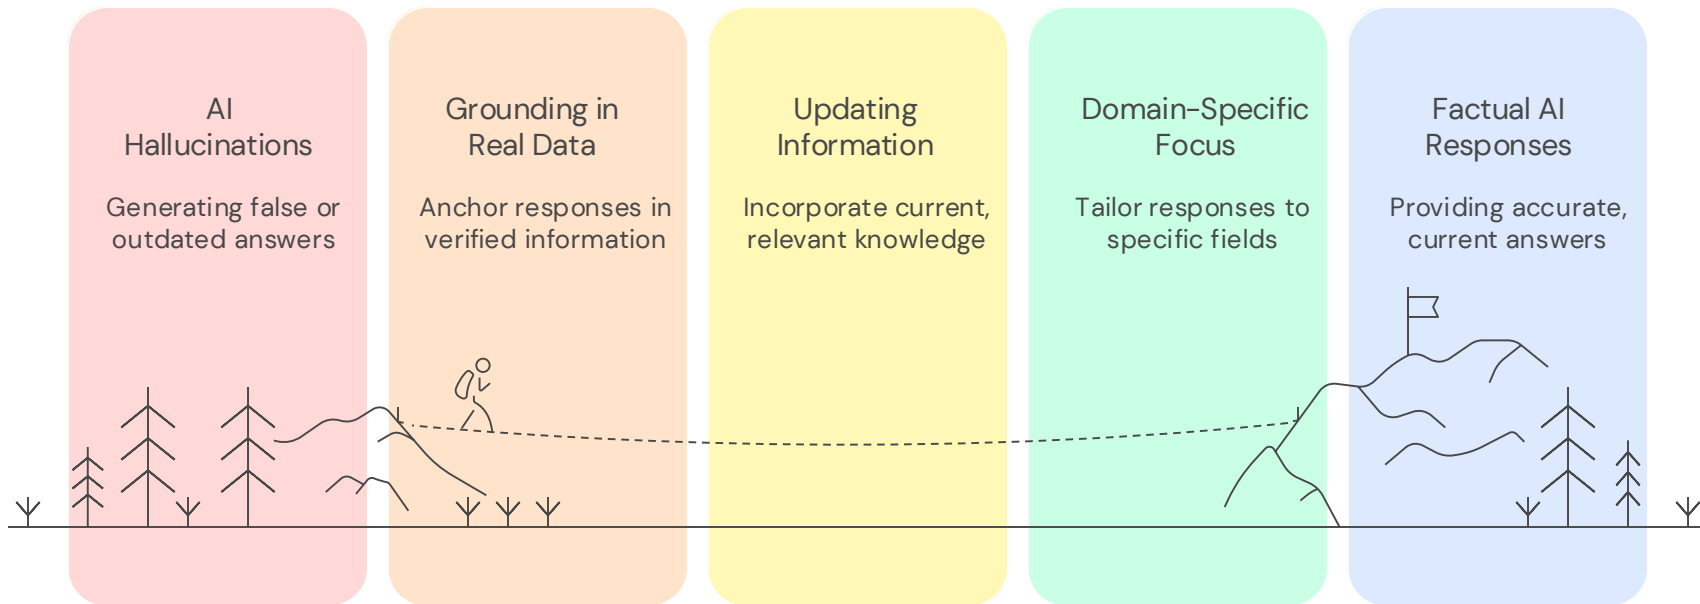
### Medical/Legal Q&A

Surfacing vetted references from trusted databases to keep outputs accurate and cite-able.

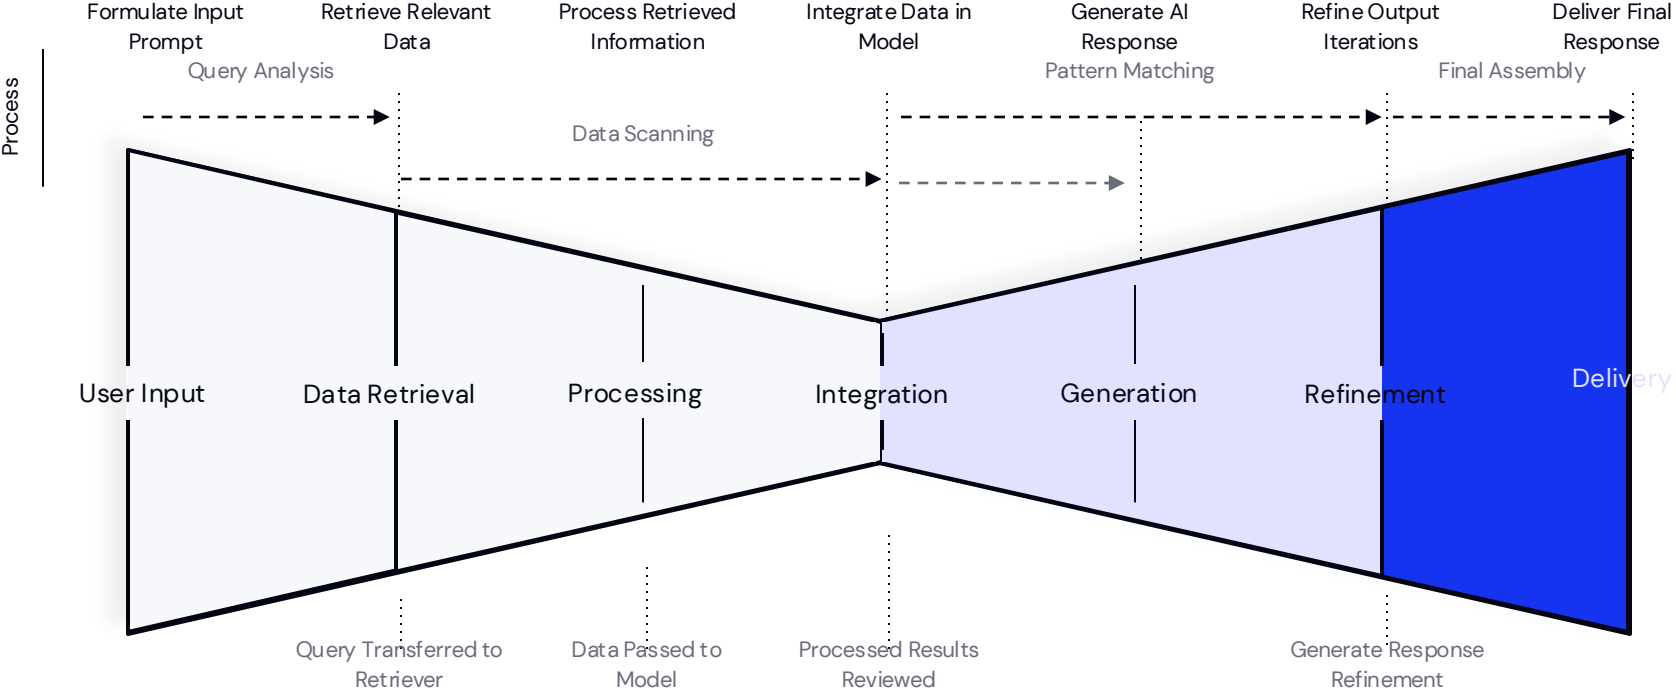
# Optimizing LLMs with RAG



# Reliable Answers with RAG



# RAG Pipeline Visual Guide



# Agentic AI Orchestration

High-level overview of Agentic AI Orchestration for LLMs.

01

## Implement an Orchestration Framework

- Wrap the base LLM with a control layer (E.g., CrewAI, LangGraph, AutoGen, LlamaIndex, etc.)
- Enables planning, task decomposition, and coordination of multi-step workflows

02

## Integrate Tools, Memory, and Context

- Connect the agent(s) to external tools (APIs, code execution, retrieval systems)
- Provide memory (short- and long-term) so it can carry context across steps

03

## Implement Reasoning & Safety Loops

- Add logic for reflection, evaluation, and self-correction
- Enforce guardrails, moderation, and approval gates to control behavior and risks

### Agentic AI Core Components



# Agentic AI – Analysis

## Benefits

Agentic AI lets LLMs handle complex, multi-step tasks autonomously, greatly boosting performance and capability without changing their core weights.

## Drawbacks

It adds architectural complexity, increases latency and cost, and requires strong safety mechanisms to prevent errors from compounding.

## Sample Use Cases

### Research Agents

Orchestrate multi-step searching, reading, and summarizing across sources.

### Workflow Automation

Chain tools and memory to automate structured business processes.

### Coding Agents

Combine planning, tool use, and feedback loops to iteratively write and debug code.

# Key Traits of Agentic AI



## Goal-Oriented

Understands objectives.



## Autonomous

Executes tasks without input.



## Reasoning

Breaks down and prioritizes.



## Self-Reflective

Continuously improves results.

Agentic AI systems prioritize objectives, autonomously execute plans, adapt reasoning to dynamic tasks, and self-reflect to optimize their outputs and actions.



# Key Insights into AI Orchestration Frameworks



## Understanding AI Orchestration

AI orchestration frameworks coordinate how multiple models, tools, and data sources work together to complete complex tasks.



## Framework Components

They typically include planning logic, memory modules, tool connectors, and control flows that manage task execution.



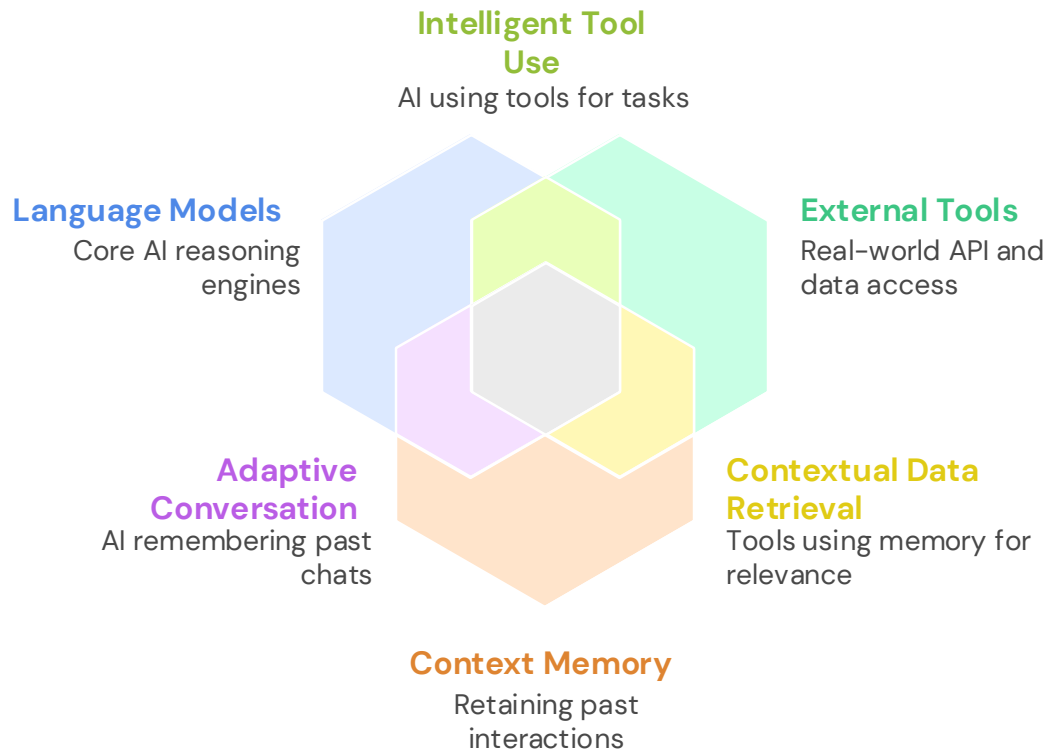
## Integration Strategies

Effective orchestration depends on clean APIs, modular design, and context-passing mechanisms to seamlessly connect systems.

## Implementation Considerations

For successful implementation of AI orchestration frameworks, ensure they are scalable, flexible, and integrate seamlessly into existing systems.

Prioritize sustainability and adaptability to dynamic changes in AI technology.



# Agentic Orchestration Frameworks

How to turn an LLM into a reasoning Agent

# From Prompt to Action Plan

Define objectives for the prompt.



Allocate resources effectively.

Identify tasks for execution.



Execute tasks as planned.

Create workflows for tasks.



Review and adapt performance.

Agentic AI Task Execution

# Autonomous Task Execution



# From LLM to Agent: The Agentic AI Design Canvas

## Agent Roles



Define the roles or personas your agents will take on (researcher, planner, coder, analyst, assistant).

## Core Capabilities



List what your agentic system can actually do (multi-step reasoning, tool use, retrieval, planning, autonomous task execution).

## Feedback & Alignment



Describe how agents receive feedback, refine behavior, and stay aligned with goals and policies.

## Model & Data Stack



Outline the components your agents rely on: LLM backbone, vector DBs, APIs, memory stores, orchestration framework, compute resources.

## Interfaces & Surfaces



Specify how users will access the agents (chat UI, voice, web app, IDE plugin, Slack bot, API).

## Agentic Behaviors



Map the core actions the system performs: planning, reasoning, task decomposition, tool invocation, error recovery, self-reflection.

## Integrations & Dependencies



List external systems and services the agents must connect to: APIs, databases, CRMs, code execution environments, monitoring tools.

## Metrics & Outcomes



Define success measures: task success rate, latency, user satisfaction, cost per run, autonomy level.

## Constraints & Limits



Note system constraints and costs: token usage, API call limits, latency budgets, safety guardrails, governance policies.

# Agentic AI Challenges

## Current Challenges

1

**Unclear purpose:** It's hard to see where agentic AI fits into daily work.

2

**Complex systems:** Multi-step agents feel opaque and hard to trust.

3

**Limited readiness:** Teams lack skills and processes to support agent-driven work.

## Negative Impacts



**Slow adoption:** Humans hesitate to use tools they don't fully understand.



**Frequent errors:** Poorly configured agents produce unreliable results.



**Workflow friction:** Agents don't yet fit smoothly into existing systems.



**User skepticism:** Confusion and mistakes reduce confidence in the tech.



**Operational drag:** Teams spend more time managing the tools than benefiting from them.



**Missed opportunities:** Early-stage failures make leaders hesitant to invest further.

# Agentic AI Benefits

## Future State

1

**Smarter workflows:** Agents handle routine steps so teams can focus on critical work.

2

**Seamless tools:** Agents connect data, apps, and systems behind the scenes.

3

**Adaptive support:** Agents learn from feedback to get more accurate over time.

## Positive Outcomes



**Less busywork:** Humans spend less time on repetitive, manual tasks.



**Faster progress:** Projects move quicker when agents automate prep work.



**More confidence:** Reliable outputs build trust and reduce second-guessing.



**Happier teams:** Less grind boosts morale and engagement.



**Stronger performance:** Teams can deliver higher-quality results faster.



**Space to innovate:** Freed-up time fuels creativity, relationship-building and new ideas.

# Thank you!

Questions?

